

Public AI is not a technology stack. It is an operating model.

Jaanus Väärt · AI and Data Strategy Coordinator, Ministry of Justice and Digital Affairs, Republic of Estonia
· September 2026

A practitioner's perspective. Facts and examples are drawn from published sources. The analysis is the author's own and does not represent an approved position of the Estonian government.

Estonia is often presented as a natural home for public-sector AI. The country has digital identity, interoperable registries, secure data exchange and long experience of designing services around authoritative data. These foundations make experimentation easier.

They do not answer the harder question: **when does an AI pilot become a durable public capability?**

A government can fund a prototype, procure a model or buy computing capacity without creating anything that is actually public. A capability becomes public only when access, accountability and continuity survive the initial project team, supplier and budget cycle.

Estonia's experience suggests that public AI is now more of an **operating-model problem** than a technology problem.

The limits of the success narrative

Estonia's decentralised digital state gives ministries and agencies substantial responsibility for their own data, legal basis, processes and services. This is useful for experimentation: institutions close to a problem can test solutions without waiting for a single central programme to design everything.

The same model gets harder when a solution has to scale across government. Shared components can attract many participants but no clear owner. A pilot budget may cover development but not monitoring, security updates, retraining, integration, incident response or long-term service management. Responsibility can get spread thin just when the service becomes important.

What Estonia already has in place

A few published examples show both the foundations and the remaining gap:

- the national [Data and AI White Paper 2024–2030](#) sets out the direction for data and AI;
- [Bürokratt](#), a virtual assistant platform developed by the Information System Authority, is used by several state agencies and local authorities, and its code is openly published;
- [ETAIS](#) provides part of Estonia's shared high-performance computing base;
- the public catalogue on [kratid.ee](#) shows AI experiments spread across many institutions;
- [AI Leap \(TI-Hüpe\)](#) brings AI tools, including commercial ones, into upper-secondary schools through a foundation working with the Ministry of Education and Research and private-sector partners.

So the direction, the use cases, the compute and the partnerships are there. Bürokratt shows what a shared component looks like when it has a clear owner. What is still being worked out is how to make that the norm: who owns a capability once it matters, who pays to run it after the pilot, which controls apply, and how the state gets out if a supplier or platform stops fitting.

Public control does not require owning every layer

For a small state, full-stack self-sufficiency would be neither realistic nor desirable. Public control should instead be defined through three layers.

1. What must remain under national control

The state cannot outsource public authority. It must retain:

- legal responsibility for the service and its decisions;
- control over sensitive data, identity and access;
- authority to set risk thresholds and accept residual risk;
- auditability, incident management and operational continuity;
- the ability to stop, replace or redesign a system when public interests require it.

None of these require owning the hardware.

2. What can be pooled at European level

Some capabilities are expensive, scarce or more effective at scale. European cooperation can provide:

- high-performance computing and specialised support;
- stronger procurement leverage;
- shared standards, testing methods and reusable components;
- resilience through a broader infrastructure base;
- access for smaller administrations that could not build comparable capacity alone.

EuroHPC AI Factories already provide computing time and support services to European users, including public authorities. The AI Gigafactories call is a further possible capacity track, not an awarded project or settled operating model. European infrastructure can reduce dependency and improve access. But capacity only becomes public AI once someone has sorted out who gets access, who is accountable and who keeps the service running after the project ends.

3. What can reasonably be sourced commercially

Commercial providers will remain essential for models, cloud services, developer tools and specialist capabilities. The relevant question is not whether a supplier is private. It is whether dependence is governable.

Commercial provision is compatible with public AI when government can verify performance, protect data, understand costs, keep systems portable and actually exit if needed. Accountability cannot be outsourced even when infrastructure is.

A practical test for public AI

The Public AI Network describes public AI through **public access, public accountability and permanent public goods**. Governments can turn those principles into five operating questions.

Test	Question
Access	Who can use the capability, under what conditions, and how are scarce resources allocated?
Decision rights	Who can approve, change, pause or retire the service, and who answers when it fails?
Lifecycle	Who funds and operates security, monitoring, maintenance, evaluation and incident response after the pilot?
Control	Can the state audit the system, protect its data, enforce service requirements and understand the chain of responsibility?
Exit	Can data, models, interfaces and operational knowledge move to another supplier or environment without losing the public service?

If these questions have no clear answers, the state may have an AI project but not yet a public AI capability.

Open source helps, but does not operate the service

Open source can make systems easier to inspect, reuse and move. It can reduce some forms of lock-in and let public institutions build on common components.

But having the source code does not create accountability by itself. Someone still has to secure, deploy, monitor, maintain and fund the service. Open code without an owner can be as unsustainable as a closed product without an exit route.

Sovereignty should be risk-based

“Sovereignty” is sometimes framed as a binary choice: build everything nationally or accept dependence on global suppliers. A more useful definition is **the ability to make and enforce risk-based choices**. Where the servers sit is only part of the test. The bigger part is whether public institutions can actually exercise control, meet their legal obligations and keep critical services running.

Different workloads need different levels of control. General productivity tools, ordinary administrative services, sensitive personal-data processing and national-security workloads should not automatically use the same infrastructure or contract model. Architecture should follow workload, legal authority, threat model and continuity needs, not slogans or someone's favourite technology.

For Estonia, the practical direction is layered: keep national control where public authority, sensitive data or continuity demand it; pool expensive capabilities where European cooperation creates leverage; and use commercial supply where dependencies stay visible and the state can still switch or exit.

What others should copy, and what they shouldn't

Other governments should not start by copying Estonia's visible application layer. The more transferable lesson is to build the operating foundations:

- authoritative and governable data;
- interoperability and reusable components;
- accountable service ownership;
- procurement capability that turns public values into enforceable requirements;
- logging, monitoring and human oversight;
- institutional memory and credible exit options.

Estonia's digital foundations mean it hits the next problem earlier than most. Using AI in government is no longer the hard part. The hard part is building on a mix of public and private technology and still being able to make the key decisions, meet legal obligations and keep the service running when a supplier, budget or team changes.

Who built the model matters less than that.